

# **Building Bazel Packages for AI/ML: SciPy, PyTorch, and Beyond**

Conference Proposal

PyData Seattle 2025

## **Speakers:**

Ramesh Oswal

Jiten Oswal

November 8, 2025

# Contents

<b>1</b>	<b>Proposal Overview</b>	<b>3</b>
1.1	Title . . . . .	3
1.2	Speakers . . . . .	3
1.3	Session Duration . . . . .	3
1.4	Target Audience . . . . .	3
<b>2</b>	<b>Abstract</b>	<b>3</b>
<b>3</b>	<b>Session Outline</b>	<b>3</b>
3.1	Introduction to Bazel for AI/ML (20 minutes) . . . . .	3
3.2	Challenges in AI/ML Bazel Packaging (30 minutes) . . . . .	4
3.3	Strategies for Packaging (30 minutes) . . . . .	4
3.4	Best Practices for Distribution and Maintenance (20 minutes) . . . . .	5
3.5	Hands-on Demo (144 minutes) . . . . .	5
3.6	Q&A and Open Discussion (36 minutes) . . . . .	6
<b>4</b>	<b>Key Takeaways</b>	<b>6</b>
<b>5</b>	<b>Prerequisites</b>	<b>6</b>
<b>6</b>	<b>Materials Provided</b>	<b>7</b>
<b>7</b>	<b>Speaker Information</b>	<b>7</b>
7.1	Ramesh Oswal . . . . .	7
7.2	Jiten Oswal . . . . .	7
<b>8</b>	<b>Why This Session Matters</b>	<b>7</b>
<b>9</b>	<b>Expected Outcomes</b>	<b>8</b>
<b>10</b>	<b>Contact Information</b>	<b>8</b>

# 1 Proposal Overview

## 1.1 Title

Building Bazel Packages for AI/ML: SciPy, PyTorch, and Beyond

## 1.2 Speakers

- **Ramesh Oswal**
- **Jiten Oswal**

## 1.3 Session Duration

280 minutes (4 hours 40 minutes)

## 1.4 Target Audience

- Software engineers working with AI/ML infrastructure
- DevOps and MLOps practitioners
- Data scientists interested in reproducible builds
- Build system engineers
- Researchers managing complex dependency chains

# 2 Abstract

AI/ML workloads depend heavily on complex software stacks, including numerical computing libraries (SciPy, NumPy), deep learning frameworks (PyTorch, TensorFlow), and specialized toolchains (CUDA, cuDNN). However, integrating these dependencies into Bazel-based workflows remains challenging due to compatibility issues, dependency resolution, and performance optimization.

This session explores the process of creating and maintaining Bazel packages for key AI/ML libraries, ensuring reproducibility, performance, and ease of use for researchers and engineers.

# 3 Session Outline

## 3.1 Introduction to Bazel for AI/ML (20 minutes)

This section introduces participants to Bazel and its relevance for AI/ML workflows:

- Overview of Bazel build system fundamentals
- Why Bazel matters for AI/ML workloads
- Benefits of Bazel for dependency management

- Reproducibility and hermetic builds
- Cross-platform compatibility considerations
- Comparison with traditional Python packaging (pip, conda)

**Learning Outcomes:**

- Understand Bazel's core concepts and terminology
- Recognize the advantages of Bazel for AI/ML projects
- Identify use cases where Bazel adds value

### 3.2 Challenges in AI/ML Bazel Packaging (30 minutes)

An in-depth exploration of the unique challenges encountered when packaging AI/ML libraries with Bazel:

- **Transitive Dependencies:** Managing complex dependency graphs
- **Build System Differences:** Bridging CMake, setuptools, and Bazel
- **GPU Acceleration:** Integrating CUDA, cuDNN, and hardware-specific toolchains
- **Binary Compatibility:** ABI compatibility across different platforms
- **Large Binary Sizes:** Handling large ML model files and libraries
- **Version Conflicts:** Resolving conflicts between library versions
- **Native Extensions:** Compiling C/C++ extensions for Python libraries

**Learning Outcomes:**

- Identify common pitfalls in AI/ML library packaging
- Understand dependency resolution complexities
- Recognize GPU toolchain integration challenges

### 3.3 Strategies for Packaging (30 minutes)

Practical strategies and techniques for successfully packaging AI/ML libraries:

- **SciPy Stack:** Building NumPy, SciPy, and related numerical libraries
- **PyTorch:** Creating Bazel rules for PyTorch with CUDA support
- **TensorFlow:** Leveraging existing Bazel infrastructure
- **Dependency Management:** Strategies for managing transitive dependencies
- **Performance Optimization:** Ensuring optimal build and runtime performance
- **Cross-Platform Support:** Building for Linux, macOS, and Windows

- **Testing and Validation:** Ensuring package correctness and functionality

#### Learning Outcomes:

- Apply proven packaging strategies to common libraries
- Optimize builds for performance and compatibility
- Implement effective testing strategies

### 3.4 Best Practices for Distribution and Maintenance (20 minutes)

Guidelines for maintaining and distributing Bazel packages in production environments:

- **Version Management:** Semantic versioning and release strategies
- **Repository Structure:** Organizing Bazel workspaces effectively
- **Distribution Channels:** Using Bazel Central Registry and private registries
- **Documentation:** Creating user-friendly package documentation
- **CI/CD Integration:** Automating build and test pipelines
- **Community Engagement:** Contributing to open-source Bazel rules
- **Monitoring and Updates:** Keeping packages up-to-date with upstream changes

#### Learning Outcomes:

- Establish sustainable maintenance workflows
- Implement effective distribution strategies
- Integrate packages into CI/CD pipelines

### 3.5 Hands-on Demo (144 minutes)

An extensive practical demonstration where participants follow along:

- **Environment Setup:** Configuring Bazel for AI/ML development
- **Building SciPy:** Step-by-step package creation for SciPy
- **Integrating PyTorch:** Adding PyTorch with GPU support
- **Creating Custom Rules:** Writing Bazel rules for custom libraries
- **Dependency Resolution:** Handling complex dependency chains
- **Testing and Validation:** Running tests to verify package functionality
- **Performance Benchmarking:** Comparing build times and runtime performance
- **Troubleshooting:** Common issues and debugging techniques

## **Learning Outcomes:**

- Build complete Bazel packages for real AI/ML libraries
- Debug and troubleshoot packaging issues
- Apply learned techniques to custom projects

## **3.6 Q&A and Open Discussion (36 minutes)**

An interactive session for:

- Addressing participant questions
- Discussing specific use cases and challenges
- Sharing experiences and best practices
- Exploring future directions for Bazel in AI/ML
- Networking and community building

## **4 Key Takeaways**

Participants will leave this session with:

1. **Practical Knowledge:** Hands-on experience building Bazel packages for AI/ML libraries
2. **Problem-Solving Skills:** Techniques to overcome common packaging challenges
3. **Best Practices:** Industry-standard approaches to package maintenance and distribution
4. **Reproducible Workflows:** Ability to create hermetic, reproducible build environments
5. **Community Resources:** Knowledge of community tools, resources, and support channels

## **5 Prerequisites**

To get the most out of this session, participants should have:

- Basic understanding of build systems and dependency management
- Familiarity with Python and AI/ML libraries
- Basic command-line proficiency
- (Optional) Prior exposure to Bazel or similar build systems
- Laptop with Docker or Bazel installed (for hands-on portion)

## 6 Materials Provided

- Complete Bazel workspace examples
- Sample BUILD files for SciPy, PyTorch, and other libraries
- Documentation and reference guides
- Troubleshooting checklists
- Links to additional resources and community support

## 7 Speaker Information

### 7.1 Ramesh Oswal

Ramesh Oswal is an experienced software engineer specializing in build systems and AI/ML infrastructure. With extensive experience in Bazel and large-scale dependency management, Ramesh has contributed to numerous open-source projects focused on reproducible ML workflows.

### 7.2 Jiten Oswal

Jiten Oswal brings expertise in DevOps, MLOps, and infrastructure automation. His work focuses on creating efficient, scalable build pipelines for AI/ML applications, with a particular emphasis on cross-platform compatibility and performance optimization.

## 8 Why This Session Matters

As AI/ML workloads become increasingly complex and production-critical, the need for reliable, reproducible build systems has never been greater. Traditional Python packaging approaches often fall short when dealing with:

- Complex native dependencies (CUDA, cuDNN, MKL)
- Large-scale monorepo structures
- Strict reproducibility requirements
- Cross-platform deployment needs
- High-performance computing environments

This session addresses these challenges head-on, providing practical solutions that teams can implement immediately. By the end of the session, participants will have the skills and confidence to manage their own AI/ML dependencies using Bazel, leading to more reliable and maintainable systems.

## **9 Expected Outcomes**

After attending this session, participants will be able to:

1. Create Bazel packages for popular AI/ML libraries
2. Integrate GPU acceleration toolchains into Bazel builds
3. Manage complex dependency graphs effectively
4. Set up reproducible build environments for AI/ML projects
5. Troubleshoot common packaging issues
6. Contribute to the Bazel AI/ML ecosystem

## **10 Contact Information**

For questions or additional information about this proposal, please contact the speakers through the GitHub repository:

<https://github.com/RameshOswal/pydata-2025-conference-bazel-ai-pkgs>

---

*Thank you for considering this proposal for PyData Seattle 2025!*